

Chanukya Gattu

Senior Software Engineer / Distributed Systems / Backend & Real-Time Streaming Platforms

845-443-6258 | chanukyasharma@gmail.com | Dallas, TX, United States | [Portfolio](#) | [LinkedIn](#) | [GitHub](#)

SUMMARY

Senior Software Engineer with **10+ years** building distributed backend and real-time data platforms at JPMorgan Chase, Capital One, AT&T, and Verizon, specializing in Java, Spring Boot, Kafka, Flink, and AWS. At JPMorgan, own the API platform for a wealth-management trading product serving **100K+** clients and advisors, and architected Continuum, a multi-tenant Flink platform running **1,000+** tenant pipelines. Focused on correctness and reliability in distributed systems, including idempotent processing, exactly-once semantics, and schema evolution.

WORK EXPERIENCE

Senior Software Engineer

JP Morgan Chase

Aug '22 — Present

Plano, United States

- **Digital Wealth Management Trading Platform** | *Java, Spring Boot, Kafka, AWS ECS, Terraform, Datadog*
 - Own the backend API layer for a digital wealth-management trading platform used daily by **100K+** clients and financial advisors, serving real-time portfolio reads at **sub-200ms** latency.
 - Architected Spring Boot microservices across portfolio, goals, and investment domains, maintaining consistency between synchronous REST reads and asynchronous Kafka-driven investment triggers and goal-state transitions.
 - Guaranteed financial data correctness under retries and replays with idempotent Kafka consumers, explicit retry contracts, backpressure handling, and schema-versioned events.
 - Automated infrastructure and deployments with Terraform, Docker, and AWS ECS, making every service independently deployable and cutting environment setup and release cycle time by **40%**.
 - Defined backend guardrails (versioned API contracts, service boundaries, RBAC, production-readiness checks) with product, platform, and security teams, adopted across **10+** services; lead design reviews and mentor **3** engineers.
- **Continuum: Multi-Tenant Streaming Data Platform** | *Apache Flink, Kafka, AWS EKS, S3, Glue, Athena, Groovy*
 - Architected and delivered Continuum, a multi-tenant streaming data platform on AWS EKS running **1,000+** isolated tenant pipelines and processing millions of events per day.
 - Designed a Groovy-based declarative DSL that lets data engineering teams define pipelines as composable DAGs without writing Flink code, cutting pipeline onboarding from **weeks to hours**.
 - Built the Apache Flink execution engine: each tenant pipeline runs as an isolated job with **exactly-once** semantics, checkpoint-triggered Parquet writes to S3, and Glue Catalog registration for Athena queries.
 - Improved pipeline reliability with stateful processing, checkpointing, and automated failure recovery; eliminated Glue API throttling by batching partition registration.
 - Enforced tenant isolation via S3 prefixes, IAM policies; autoscaled EKS pods on per-tenant Kafka consumer lag.

Senior Software Engineer

AT&T

Aug '21 — Aug '22

Plano, United States

- **Cloud Migration - BILL PAY** | *Java, Spring Boot, Kafka, AWS, Terraform, Datadog*
 - Led migration of a legacy on-prem platform to AWS (ECS and Kubernetes) with no disruption to live operations, moving **15+** services onto Terraform-managed infrastructure.
 - Decomposed a monolith, where every release touched every component, into independently deployable Spring Boot microservices with clean REST contracts and per-service scaling.
 - Built Jenkins and GitHub Actions CI/CD pipelines with Terraform provisioning, enabling independent deployments and reducing release effort by **30%**.
 - Built observability with Datadog, CloudWatch, and Grafana dashboards and alerting, speeding incident detection and improving MTTR by **35%**.

Senior Software Engineer

Capital One

Jan '20 — Aug '21

Plano, United States

- **Auto Loans Backend Platform** | *Java, Spring Boot, Kafka, AWS, Terraform, Datadog*
 - Built Java Spring Boot microservices for Auto Loans origination, servicing, and underwriting APIs, keeping partner-facing REST APIs responsive under peak traffic.
 - Built Kafka pipelines for loan, payment, and underwriting events with strict ordering and delivery guarantees.
 - Raised event and batch throughput by **35%** under peak load using Java concurrency APIs and multithreading.
 - Secured PCI-scoped data with field-level encryption, tokenization, and RBAC across all API surfaces.

Software Engineer

Verizon

May '15 — Jan '20

Irving, United States

- **OPTIX Ordering Suite** | *Java, Spring, RESTFUL API's, PCF, PostgreSQL*
 - Built RESTFUL APIs on Spring framework and integration modules for the OPTIX Ordering Suite, supporting FiOS order management and provisioning for thousands of call-center operators daily.
 - Reduced order-processing latency by **~25%** through SQL query tuning, caching, and asynchronous processing.
 - Migrated **20+** legacy SOAP endpoints to RESTful APIs, reducing service coupling across ordering integrations.

OPEN-SOURCE PROJECTS

StreamForge [Link](#)

- **YAML-to-Flink pipeline compiler**
 - Compiles source | transform | sink YAML into Flink jobs with exactly-once semantics, resolving connectors and schema, and registering output in the Glue catalog. Java 17, Flink 1.18+, Docker.
 - Designed a pluggable connector SPI for Kafka, S3, Iceberg, and Snowflake, so new sinks are added as modules rather than forks; includes a control plane, checkpoint-driven emission, and per-tenant isolation.

Flink for Dummies [Link](#)

- **Interactive guide to the Flink runtime**
 - Four browser labs that port Flink's real key-group hashing (murmurHash, KeyGroupRangeAssignment), watermark rules, aligned and unaligned checkpoint barriers, and backpressure metrics. TypeScript, Docusaurus.
 - 60-page textbook from events to runtime internals; each page covers the problem, mental model, what Flink does internally, how it fails, and what changes at scale.

SKILLS

Languages & Frameworks Java/J2EE, Hibernate, Spring Boot, Python, SQL, Pyspark, Mockito, Shell

Architecture & Design Microservices, RESTful API design, Event-Driven Systems, 12-Factor App Principles, Distributed Systems

Messaging & Data Streaming Apache Kafka, Flink, RabbitMQ, Redis, PostgreSQL, Cassandra, AWS Aurora

Cloud & Infrastructure AWS (Glue, IAM, Lambda, ECS, EKS, S3, RDS, API Gateway), Docker, Kubernetes, Terraform, Pivotal Cloud Foundry (PCF)

Monitoring & Logging Datadog, AWS CloudWatch, Prometheus, Dynatrace, Splunk, Grafana, ELK

DevOps & CI/CD Jenkins, GitHub Actions, Maven, Spinnaker

CERTIFICATIONS

[AWS Certified Solutions Architect - Associate SAA-C03](#)

[HashiCorp Certified: Terraform Associate \(004\)](#)

[AWS Certified Developer Associate DVA-C02](#)

[AWS Certified Cloud Practitioner](#)

EDUCATION

Master of Science in Electrical Engineering, State University of New York Aug '13 — May '15
New Paltz, New York, United States

- Specialized in embedded systems, algorithms, distributed computing, networking, and database management. Gained hands-on experience through Java-based projects and IoT applications, bridging hardware and software development.

Bachelor of Technology in Electrical Engineering, Kakatiya University Aug '09 — May '13
Warangal, Telangana, India

- Built a solid foundation in electrical systems, programming, and applied mathematics, with a strong emphasis on analytical problem solving and system modeling.